

Summer Class: Machine Learning and Data Science

Shuaifeng Li

Department of Physics, University of Michigan

June & July 2025

Table of Contents

- 1 Overview of Machine Learning
- 2 Convex Optimization
- 3 Singular Value Decomposition**
- 4 Fourier Analysis
- 5 Linear Regression
- 6 Dynamical Systems

Matrix as a Vector Transformer

Any matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ transforms a vector $\mathbf{x} \in \mathbb{R}^n$ into a new vector $\mathbf{y} \in \mathbb{R}^m$:

$$\mathbf{y} = \mathbf{A}\mathbf{x}$$

Matrix as a Vector Transformer

Any matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ transforms a vector $\mathbf{x} \in \mathbb{R}^n$ into a new vector $\mathbf{y} \in \mathbb{R}^m$:

$$\mathbf{y} = \mathbf{A}\mathbf{x}$$

Example:

$$\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 0 & 1 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \Rightarrow \mathbf{y} = \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$

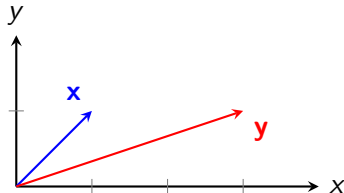
Matrix as a Vector Transformer

Any matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ transforms a vector $\mathbf{x} \in \mathbb{R}^n$ into a new vector $\mathbf{y} \in \mathbb{R}^m$:

$$\mathbf{y} = \mathbf{A}\mathbf{x}$$

Example:

$$\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 0 & 1 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \Rightarrow \mathbf{y} = \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$



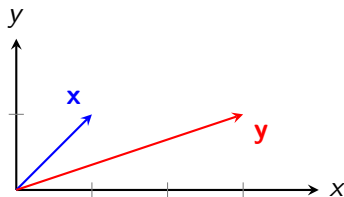
Matrix as a Vector Transformer

Any matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ transforms a vector $\mathbf{x} \in \mathbb{R}^n$ into a new vector $\mathbf{y} \in \mathbb{R}^m$:

$$\mathbf{y} = \mathbf{A}\mathbf{x}$$

Example:

$$\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 0 & 1 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \Rightarrow \mathbf{y} = \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$



What does this mean geometrically?

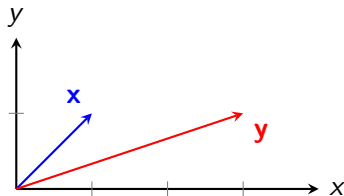
Matrix as a Vector Transformer

Any matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ transforms a vector $\mathbf{x} \in \mathbb{R}^n$ into a new vector $\mathbf{y} \in \mathbb{R}^m$:

$$\mathbf{y} = \mathbf{A}\mathbf{x}$$

Example:

$$\mathbf{A} = \begin{bmatrix} 2 & 1 \\ 0 & 1 \end{bmatrix}, \quad \mathbf{x} = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \Rightarrow \mathbf{y} = \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$

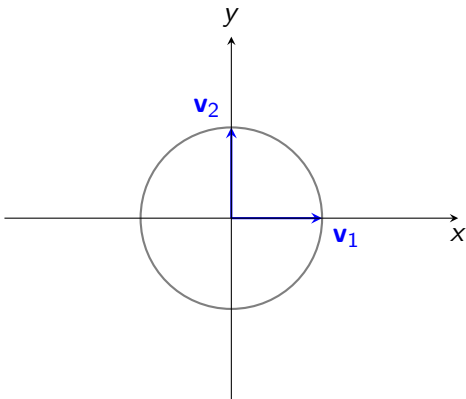


What does this mean geometrically?

- Changes length (scaling).
- Changes direction (rotation/skewing).
- May collapse dimensions.

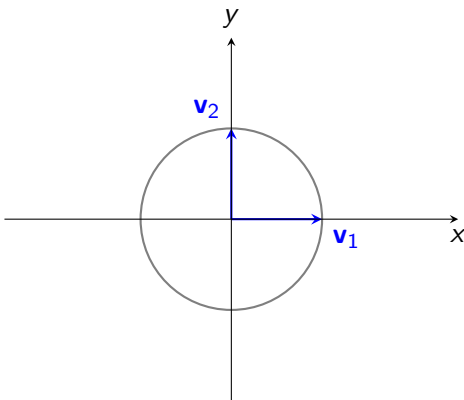
Transforming a Set of Vectors

Input space: Unit Circle

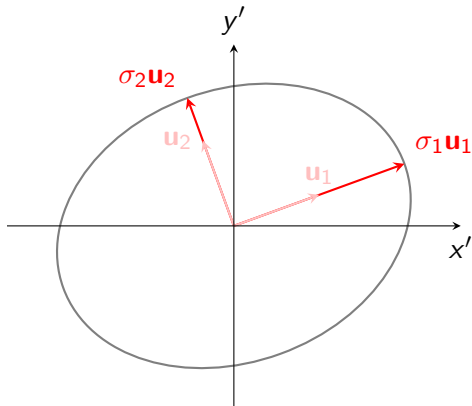


Transforming a Set of Vectors

Input space: Unit Circle

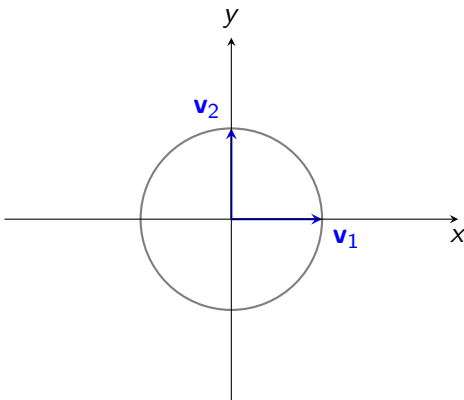


After A , Output space: Ellipse

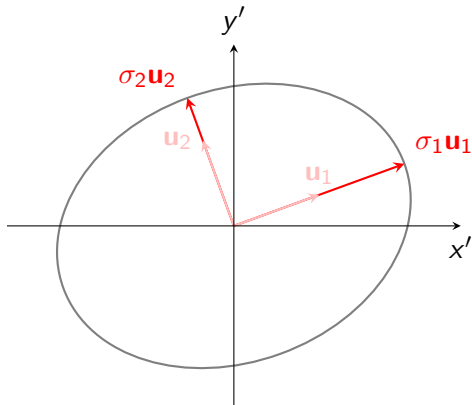


Transforming a Set of Vectors

Input space: Unit Circle



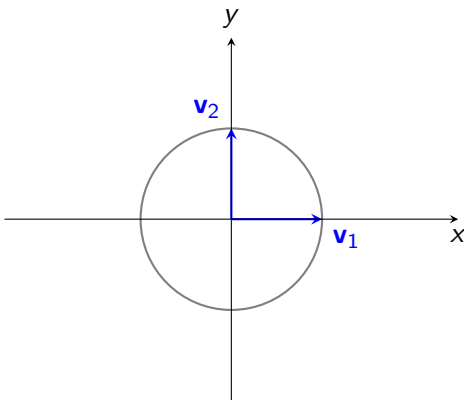
After A , Output space: Ellipse



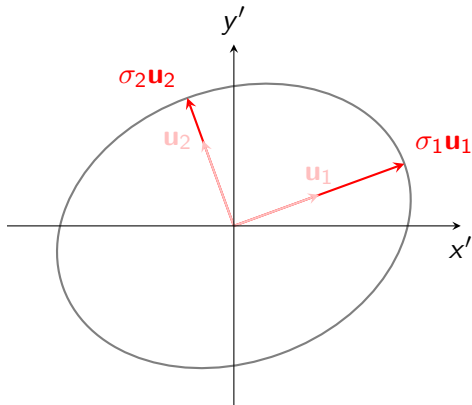
A matrix maps unit vectors along $\mathbf{v}_1, \mathbf{v}_2$ to scaled orthonormal vectors along $\sigma_1 \mathbf{u}_1, \sigma_2 \mathbf{u}_2$:

Transforming a Set of Vectors

Input space: Unit Circle



After \mathbf{A} , Output space: Ellipse



\mathbf{A} matrix maps unit vectors along $\mathbf{v}_1, \mathbf{v}_2$ to scaled orthonormal vectors along $\sigma_1 \mathbf{u}_1, \sigma_2 \mathbf{u}_2$:

$$\mathbf{A} \mathbf{v}_i = \sigma_i \mathbf{u}_i \Rightarrow \mathbf{A} \mathbf{V} = \mathbf{U} \mathbf{\Sigma} \Rightarrow \mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^*.$$

What is Singular Value Decomposition?

- The Singular Value Decomposition (SVD) is a fundamental matrix factorization in linear algebra. It asserts that any real or complex matrix $A \in \mathbb{C}^{m \times n}$ can be decomposed into the product of three specific matrices.

$$A = U\Sigma V^*$$

What is Singular Value Decomposition?

- The Singular Value Decomposition (SVD) is a fundamental matrix factorization in linear algebra. It asserts that any real or complex matrix $A \in \mathbb{C}^{m \times n}$ can be decomposed into the product of three specific matrices.

$$A = U\Sigma V^*$$

- $U \in \mathbb{C}^{m \times m}$: An orthogonal (or unitary for complex) matrix whose columns are the left singular vectors.
- $\Sigma \in \mathbb{C}^{m \times n}$: A diagonal matrix (with non-negative real numbers in decreasing order on the diagonal) containing the singular values σ .
- $V \in \mathbb{C}^{n \times n}$: An orthogonal (or unitary for complex) matrix whose columns are the right singular vectors.

Properties of SVD Components

- Orthogonality of U and V :

Properties of SVD Components

- Orthogonality of U and V :
 - $U^T U = I_m$ and $V^T V = I_n$.

Properties of SVD Components

- Orthogonality of U and V :
 - $U^T U = I_m$ and $V^T V = I_n$.
 - The columns of U/V are mutually orthogonal unit vectors.
 - This orthogonality is crucial as it implies that U and V represent rotations and/or reflections, which preserve lengths in a geometric sense.

Properties of SVD Components

- Orthogonality of U and V :
 - $U^T U = I_m$ and $V^T V = I_n$.
 - The columns of U/V are mutually orthogonal unit vectors.
 - This orthogonality is crucial as it implies that U and V represent rotations and/or reflections, which preserve lengths in a geometric sense.
- Ordering of Singular Values:

Properties of SVD Components

- Orthogonality of U and V :
 - $U^T U = I_m$ and $V^T V = I_n$.
 - The columns of U/V are mutually orthogonal unit vectors.
 - This orthogonality is crucial as it implies that U and V represent rotations and/or reflections, which preserve lengths in a geometric sense.
- Ordering of Singular Values:
 - $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, where $r = \text{rank}(A)$. Any remaining diagonal entries are zero.

Properties of SVD Components

- Orthogonality of U and V :
 - $U^T U = I_m$ and $V^T V = I_n$.
 - The columns of U/V are mutually orthogonal unit vectors.
 - This orthogonality is crucial as it implies that U and V represent rotations and/or reflections, which preserve lengths in a geometric sense.
- Ordering of Singular Values:
 - $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, where $r = \text{rank}(A)$. Any remaining diagonal entries are zero.
 - This ordering is fundamental for low-rank approximation and identifying dominant components. The largest singular values correspond to the directions of greatest variance or importance in the data.

Properties of SVD Components

- Orthogonality of U and V :
 - $U^T U = I_m$ and $V^T V = I_n$.
 - The columns of U/V are mutually orthogonal unit vectors.
 - This orthogonality is crucial as it implies that U and V represent rotations and/or reflections, which preserve lengths in a geometric sense.
- Ordering of Singular Values:
 - $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$, where $r = \text{rank}(A)$. Any remaining diagonal entries are zero.
 - This ordering is fundamental for low-rank approximation and identifying dominant components. The largest singular values correspond to the directions of greatest variance or importance in the data.
 - If $m > n$, $\Sigma = \begin{bmatrix} \hat{\Sigma} \in \mathbb{C}^{n \times n} \\ 0 \end{bmatrix}$. If $m < n$, $\Sigma = [\hat{\Sigma} \in \mathbb{C}^{m \times m}, 0]$. If $m = n$, $\Sigma = \hat{\Sigma} \in \mathbb{C}^{m \times n}$.

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

- 1 Rotation by V^T :

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

- 1 Rotation by V^T :
 - V^T is an orthogonal matrix. It rotates the input vectors in the domain (input space \mathbb{R}^n).
 - It aligns the basis vectors of the domain along the principal axes (directions of v_i).

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

- ① Rotation by V^T :
 - V^T is an orthogonal matrix. It rotates the input vectors in the domain (input space \mathbb{R}^n).
 - It aligns the basis vectors of the domain along the principal axes (directions of v_i).
- ② Scaling by Σ :

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

① Rotation by V^T :

- V^T is an orthogonal matrix. It rotates the input vectors in the domain (input space \mathbb{R}^n).
- It aligns the basis vectors of the domain along the principal axes (directions of v_i).

② Scaling by Σ :

- Σ is a diagonal matrix. It scales the rotated vectors along these new axes.
- The scaling factor along each axis i is the singular value σ_i . Some axes might be stretched ($\sigma_i > 1$), some shrunk ($\sigma_i < 1$), and some collapsed ($\sigma_i = 0$).

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

① Rotation by V^T :

- V^T is an orthogonal matrix. It rotates the input vectors in the domain (input space \mathbb{R}^n).
- It aligns the basis vectors of the domain along the principal axes (directions of v_i).

② Scaling by Σ :

- Σ is a diagonal matrix. It scales the rotated vectors along these new axes.
- The scaling factor along each axis i is the singular value σ_i . Some axes might be stretched ($\sigma_i > 1$), some shrunk ($\sigma_i < 1$), and some collapsed ($\sigma_i = 0$).

③ Rotation by U :

Geometric Interpretation of SVD

Any linear transformation $X_{new} = AX_0$ represented by a matrix A can be decomposed into a sequence of three fundamental geometric operations:

① Rotation by V^T :

- V^T is an orthogonal matrix. It rotates the input vectors in the domain (input space \mathbb{R}^n).
- It aligns the basis vectors of the domain along the principal axes (directions of v_i).

② Scaling by Σ :

- Σ is a diagonal matrix. It scales the rotated vectors along these new axes.
- The scaling factor along each axis i is the singular value σ_i . Some axes might be stretched ($\sigma_i > 1$), some shrunk ($\sigma_i < 1$), and some collapsed ($\sigma_i = 0$).

③ Rotation by U :

- U is an orthogonal matrix. It rotates the scaled vectors into their final positions in the codomain (output space \mathbb{R}^m).
- It aligns the scaled vectors with the basis vectors of the range (u_i).

Visualizing the Transformation

Let's start with an ellipse and apply the transformation matrix:

$$A = \begin{pmatrix} 0.4330 & -0.7500 \\ 1.2500 & -0.4330 \end{pmatrix}$$

- **Rotation (V^T):** Rotation by 30° .

$$V^T = \begin{pmatrix} 0.866 & -0.5 \\ 0.5 & 0.866 \end{pmatrix} \approx R(30^\circ)$$

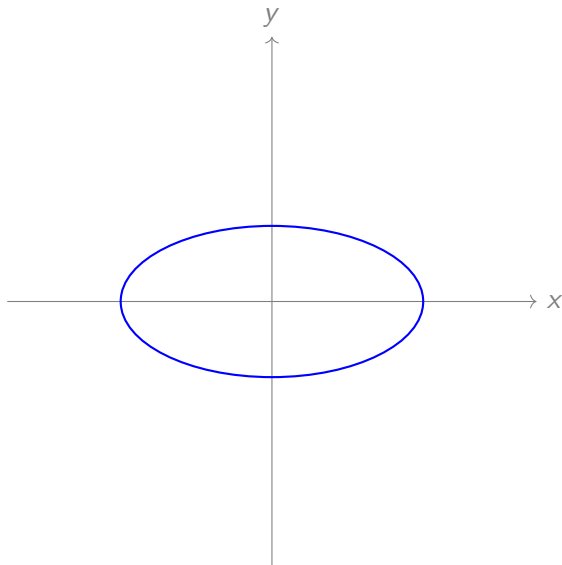
- **Scaling (Σ):**

$$\Sigma = \begin{pmatrix} 1.5 & 0 \\ 0 & 0.5 \end{pmatrix}$$

- **Rotation (U):** Rotation by 60° .

$$U = \begin{pmatrix} 0.5 & -0.866 \\ 0.866 & 0.5 \end{pmatrix} \approx R(60^\circ)$$

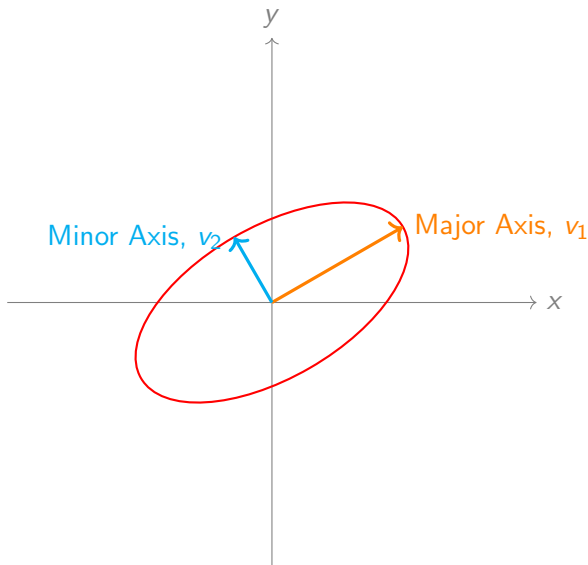
Visualizing the Transformation



Start: Ellipse

$$x_0 \in \mathbb{R}^2$$

Visualizing the Transformation

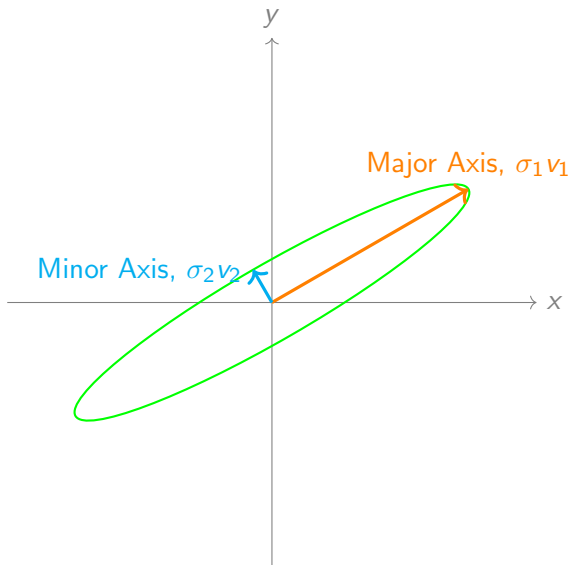


Step 1: Apply V^T

$$x_1 = V^T x_0$$

V^T rotates the entire coordinate system so that the directions we want to stretch and shrink are perfectly aligned with the main x and y axes (major and minor axes are rotated to lie along the 30° and 120° lines).

Visualizing the Transformation

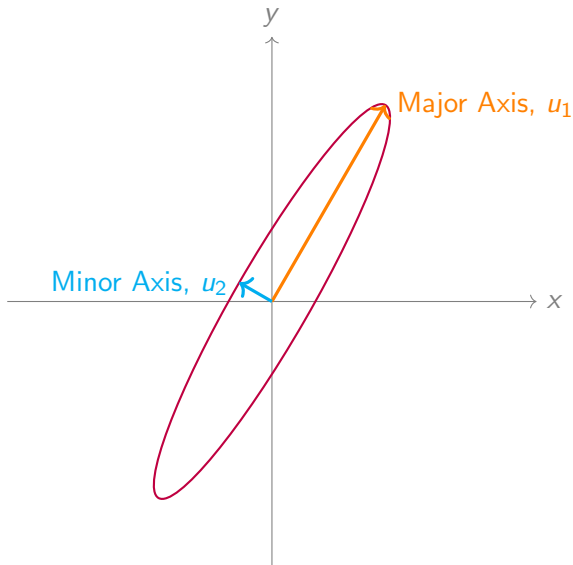


Step 2: Apply Σ

$$x_2 = \Sigma x_1 = \Sigma V^T x_0$$

It scales everything by 1.5 along the new 30° direction and by 0.5 along the new 120° direction.

Visualizing the Transformation



Step 3: Apply U

$$x_{\text{new}} = Ux_2 = U\Sigma V^T x_0$$

The columns of the U define the directions of the final ellipse's axes.

Final rotation gives:

$$x_{\text{new}} = Ax_0$$

Full vs Economy (Reduced) SVD

- Full SVD: $U \in \mathbb{C}^{m \times m}$, $\Sigma \in \mathbb{C}^{m \times n}$, $V \in \mathbb{C}^{n \times n}$.

Full vs Economy (Reduced) SVD

- Full SVD: $U \in \mathbb{C}^{m \times m}$, $\Sigma \in \mathbb{C}^{m \times n}$, $V \in \mathbb{C}^{n \times n}$.
- Economy SVD:
 - **Case 1: Tall-and-Skinny Matrix** ($m > n$)

$$A_{m \times n} = \underbrace{\begin{bmatrix} | & & | \\ u_1 & \cdots & u_n \\ | & & | \end{bmatrix}}_{U \in \mathbb{C}^{m \times n}} \underbrace{\begin{bmatrix} \sigma_1 & & 0 \\ & \ddots & \\ 0 & & \sigma_n \end{bmatrix}}_{\Sigma \in \mathbb{C}^{n \times n}} \underbrace{\begin{bmatrix} - & v_1^\top & - \\ & \vdots & \\ - & v_n^\top & - \end{bmatrix}}_{V^\top \in \mathbb{C}^{n \times n}}$$

Full vs Economy (Reduced) SVD

- Full SVD: $U \in \mathbb{C}^{m \times m}$, $\Sigma \in \mathbb{C}^{m \times n}$, $V \in \mathbb{C}^{n \times n}$.
- Economy SVD:
 - **Case 1: Tall-and-Skinny Matrix** ($m > n$)

$$A_{m \times n} = \underbrace{\begin{bmatrix} | & & | \\ u_1 & \cdots & u_n \\ | & & | \end{bmatrix}}_{U \in \mathbb{C}^{m \times n}} \underbrace{\begin{bmatrix} \sigma_1 & & 0 \\ & \ddots & \\ 0 & & \sigma_n \end{bmatrix}}_{\Sigma \in \mathbb{C}^{n \times n}} \underbrace{\begin{bmatrix} - & v_1^\top & - \\ & \vdots & \\ - & v_n^\top & - \end{bmatrix}}_{V^\top \in \mathbb{C}^{n \times n}}$$

- **Case 2: Short-and-Fat Matrix** ($m < n$)

$$A_{m \times n} = \underbrace{\begin{bmatrix} | & & | \\ u_1 & \cdots & u_m \\ | & & | \end{bmatrix}}_{U \in \mathbb{C}^{m \times m}} \underbrace{\begin{bmatrix} \sigma_1 & & 0 \\ & \ddots & \\ 0 & & \sigma_m \end{bmatrix}}_{\Sigma \in \mathbb{C}^{m \times m}} \underbrace{\begin{bmatrix} - & v_1^\top & - \\ & \vdots & \\ - & v_m^\top & - \end{bmatrix}}_{V^\top \in \mathbb{C}^{m \times n}}$$

How to Compute SVD

- Let $A = U\Sigma V^T$.

How to Compute SVD

- Let $A = U\Sigma V^T$.
- Then:
 - Row-space correlation matrix:

$$AA^T = U\Sigma V^T V\Sigma^T U^T = U\Sigma^2 U^T \Rightarrow U = \text{eigenvectors of } AA^T$$

- Column-space correlation matrix:

$$A^T A = V\Sigma^T U^T U\Sigma V^T = V\Sigma^2 V^T \Rightarrow V = \text{eigenvectors of } A^T A$$

How to Compute SVD

- Let $A = U\Sigma V^T$.
- Then:
 - Row-space correlation matrix:

$$AA^T = U\Sigma V^T V\Sigma^T U^T = U\Sigma^2 U^T \Rightarrow U = \text{eigenvectors of } AA^T$$

- Column-space correlation matrix:

$$A^T A = V\Sigma^T U^T U\Sigma V^T = V\Sigma^2 V^T \Rightarrow V = \text{eigenvectors of } A^T A$$

- Eigenvalues: Σ^2

Randomized SVD (rSVD)

Challenge:

- Standard SVD algorithms can be computationally expensive for very large matrices.

Randomized SVD (rSVD)

Challenge:

- Standard SVD algorithms can be computationally expensive for very large matrices.
- For “Big Data” scenarios where m and n are in the tens of thousands, millions, or more, computing the full SVD becomes impractical in terms of time and memory.

Randomized SVD (rSVD)

Challenge:

- Standard SVD algorithms can be computationally expensive for very large matrices.
- For “Big Data” scenarios where m and n are in the tens of thousands, millions, or more, computing the full SVD becomes impractical in terms of time and memory.
- However, many large real-world matrices are approximately low-rank, or we are only interested in a low-rank approximation (e.g., top k singular values/vectors).

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

- 1 Finding an Approximate Basis for the Span of the Columns of A .

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

- ① Finding an Approximate Basis for the Span of the Columns of A .
 - Goal: Construct a smaller matrix $Q \in \mathbb{R}^{m \times l}$ (where l is slightly larger than the target rank k) whose columns are orthonormal and span a subspace that captures most of column space of A .

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

- ① Finding an Approximate Basis for the Span of the Columns of A .
 - Goal: Construct a smaller matrix $Q \in \mathbb{R}^{m \times l}$ (where l is slightly larger than the target rank k) whose columns are orthonormal and span a subspace that captures most of column space of A .
 - How:
 - ① Generate a random matrix $\Omega \in \mathbb{R}^{n \times l}$.
 - ② Form the “sketch” or “sample” matrix $Y = A\Omega \in \mathbb{R}^{m \times l}$. Multiplying A by random vectors effectively samples its range.
 - ③ Compute an orthonormal basis for the columns of Y , typically via QR decomposition: $Y = QR$. The matrix $Q \in \mathbb{R}^{m \times l}$ is our approximate basis.

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

- ① Finding an Approximate Basis for the Span of the Columns of A .
 - Goal: Construct a smaller matrix $Q \in \mathbb{R}^{m \times l}$ (where l is slightly larger than the target rank k) whose columns are orthonormal and span a subspace that captures most of column space of A .
 - How:
 - ① Generate a random matrix $\Omega \in \mathbb{R}^{n \times l}$.
 - ② Form the “sketch” or “sample” matrix $Y = A\Omega \in \mathbb{R}^{m \times l}$. Multiplying A by random vectors effectively samples its range.
 - ③ Compute an orthonormal basis for the columns of Y , typically via QR decomposition: $Y = QR$. The matrix $Q \in \mathbb{R}^{m \times l}$ is our approximate basis.

The insight here is that random projections can preserve the essential geometric structure of the original matrix A , particularly its dominant components, allowing us to work with a much smaller sketch Y .

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

2 SVD on Smaller Matrix

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

2 SVD on Smaller Matrix

- Once we have $Q \in \mathbb{R}^{m \times l}$ with orthonormal columns, where $A \approx QQ^T A$:

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

② SVD on Smaller Matrix

- Once we have $Q \in \mathbb{R}^{m \times l}$ with orthonormal columns, where $A \approx QQ^T A$:
 - ① Form a much smaller matrix $B = Q^T A$. Dimensions of B are $l \times n$.

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

② SVD on Smaller Matrix

- Once we have $Q \in \mathbb{R}^{m \times l}$ with orthonormal columns, where $A \approx QQ^T A$:
 - ① Form a much smaller matrix $B = Q^T A$. Dimensions of B are $l \times n$.

The key is that computing SVD of B is much faster than SVD of A because $l \ll m, n$.

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

② SVD on Smaller Matrix

- Once we have $Q \in \mathbb{R}^{m \times l}$ with orthonormal columns, where $A \approx QQ^T A$:
 - ① Form a much smaller matrix $B = Q^T A$. Dimensions of B are $l \times n$.
 - ② Compute the SVD of this smaller matrix $B = \tilde{U}\Sigma V^T$.

The key is that computing SVD of B is much faster than SVD of A because $l \ll m, n$.

Basic Idea of rSVD

Randomized SVD (rSVD) algorithms aim to compute a near-optimal low-rank approximation much faster than traditional SVD, especially for large matrices. The core idea involves two stages:

② SVD on Smaller Matrix

- Once we have $Q \in \mathbb{R}^{m \times l}$ with orthonormal columns, where $A \approx QQ^T A$:
 - ① Form a much smaller matrix $B = Q^T A$. Dimensions of B are $l \times n$.
 - ② Compute the SVD of this smaller matrix $B = \tilde{U}\Sigma V^T$.
 - ③ The final SVD approximation for A is then $A \approx (Q\tilde{U})\Sigma V^T$. Let $U_{rSVD} = Q\tilde{U}$ and then $A \approx U_{rSVD}\Sigma V^T$.

The key is that computing SVD of B is much faster than SVD of A because $l \ll m, n$.

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

- 1 **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

- 1 **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)
- 2 **Sketch the range of A :** $Y = A\Omega \in \mathbb{R}^{m \times (k+p)}$

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

- ① **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)
- ② **Sketch the range of A :** $Y = A\Omega \in \mathbb{R}^{m \times (k+p)}$
- ③ **Orthonormalize by QR Decomposition:**
 $Y = QR \Rightarrow Q \in \mathbb{R}^{m \times (k+p)}$

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

- ① **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)
- ② **Sketch the range of A :** $Y = A\Omega \in \mathbb{R}^{m \times (k+p)}$
- ③ **Orthonormalize by QR Decomposition:**
 $Y = QR \Rightarrow Q \in \mathbb{R}^{m \times (k+p)}$
- ④ **Project A into lower dimension:** $B = Q^T A \in \mathbb{R}^{(k+p) \times n}$

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

- ① **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)
- ② **Sketch the range of A :** $Y = A\Omega \in \mathbb{R}^{m \times (k+p)}$
- ③ **Orthonormalize by QR Decomposition:**
 $Y = QR \Rightarrow Q \in \mathbb{R}^{m \times (k+p)}$
- ④ **Project A into lower dimension:** $B = Q^T A \in \mathbb{R}^{(k+p) \times n}$
- ⑤ **Compute SVD on small matrix:** $B = \tilde{U}\Sigma V^T$

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

① **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)

② **Sketch the range of A :** $Y = A\Omega \in \mathbb{R}^{m \times (k+p)}$

③ **Orthonormalize by QR Decomposition:**

$$Y = QR \Rightarrow Q \in \mathbb{R}^{m \times (k+p)}$$

④ **Project A into lower dimension:** $B = Q^T A \in \mathbb{R}^{(k+p) \times n}$

⑤ **Compute SVD on small matrix:** $B = \tilde{U}\Sigma V^T$

⑥ **Form approximate SVD of A :**

$$A \approx (Q\tilde{U})\Sigma V^T \Rightarrow A_{rSVD} = A_k = U_k \Sigma_k V_k^T$$

rSVD: Step-by-Step

Approximate the top k singular values/vectors of a large matrix $A \in \mathbb{R}^{m \times n}$

① **Generate random test matrix:** $\Omega \in \mathbb{R}^{n \times (k+p)}$ ($l = k + p$)

② **Sketch the range of A :** $Y = A\Omega \in \mathbb{R}^{m \times (k+p)}$

③ **Orthonormalize by QR Decomposition:**

$$Y = QR \Rightarrow Q \in \mathbb{R}^{m \times (k+p)}$$

④ **Project A into lower dimension:** $B = Q^T A \in \mathbb{R}^{(k+p) \times n}$

⑤ **Compute SVD on small matrix:** $B = \tilde{U}\Sigma V^T$

⑥ **Form approximate SVD of A :**

$$A \approx (Q\tilde{U})\Sigma V^T \Rightarrow A_{rSVD} = A_k = U_k \Sigma_k V_k^T$$

MATLAB Demonstration.

Matrix Approximation

Matrix Approximation

SVD as a Sum of Rank-One Matrices:

Matrix Approximation

SVD as a Sum of Rank-One Matrices:

The SVD $A = U\Sigma V^T$ can be written as an outer product expansion:

$$A = \sum_{i=1}^r \sigma_i u_i v_i^T$$

Matrix Approximation

SVD as a Sum of Rank-One Matrices:

The SVD $A = U\Sigma V^T$ can be written as an outer product expansion:

$$A = \sum_{i=1}^r \sigma_i u_i v_i^T$$

where:

- $r = \text{rank}(A)$.
- σ_i is the i -th singular value.
- u_i is the i -th left singular vector (column of U).
- v_i is the i -th right singular vector (column of V).
- Each item $\sigma_i u_i v_i^T$ is an $m \times n$ matrix of rank 1.

Matrix Approximation

SVD as a Sum of Rank-One Matrices:

The SVD $A = U\Sigma V^T$ can be written as an outer product expansion:

$$A = \sum_{i=1}^r \sigma_i u_i v_i^T$$

where:

- $r = \text{rank}(A)$.
- σ_i is the i -th singular value.
- u_i is the i -th left singular vector (column of U).
- v_i is the i -th right singular vector (column of V).
- Each item $\sigma_i u_i v_i^T$ is an $m \times n$ matrix of rank 1.

This means any matrix A can be expressed as a weighted sum of r rank-one matrices. The weights are the singular values σ_i . Since $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r \geq 0$, the first few terms in this sum capture the most “significant” parts of the matrix A .

Matrix Approximation

Low-Rank Approximation:

Matrix Approximation

Low-Rank Approximation:

If we truncate the sum of rank-one matrices after k terms ($k < r$):

$$A_k = \sum_{i=1}^k \sigma_i u_i v_i^T.$$

This matrix A_k is a rank- k approximation of A .

Matrix Approximation

Low-Rank Approximation:

If we truncate the sum of rank-one matrices after k terms ($k < r$):

$$A_k = \sum_{i=1}^k \sigma_i u_i v_i^T.$$

This matrix A_k is a rank- k approximation of A .

- A_k is constructed using the k largest singular values and their corresponding singular vectors.
- This process effectively filters out information associated with smaller singular values, which often corresponds to noise or less significant details.

Optimal Truncation – Eckart-Young-Mirsky Theorem

The Eckart-Young-Mirsky theorem formalizes the optimality of the SVD-based low-rank approximation.

Optimal Truncation – Eckart-Young-Mirsky Theorem

The Eckart-Young-Mirsky theorem formalizes the optimality of the SVD-based low-rank approximation.

- Frobenius norm

$$\|A\|_F = \left(\sum_{i=1}^n \sum_{j=1}^p |a_{ij}|^2 \right)^{\frac{1}{2}} = (\text{tr} A^T A)^{\frac{1}{2}}, \quad \|A\|_F = \left(\sum_{i=1}^r \sigma_i^2 \right)^{\frac{1}{2}}$$

Optimal Truncation – Eckart-Young-Mirsky Theorem

The Eckart-Young-Mirsky theorem formalizes the optimality of the SVD-based low-rank approximation.

- Frobenius norm

$$\|A\|_F = \left(\sum_{i=1}^n \sum_{j=1}^p |a_{ij}|^2 \right)^{\frac{1}{2}} = (\text{tr} A^T A)^{\frac{1}{2}}, \quad \|A\|_F = \left(\sum_{i=1}^r \sigma_i^2 \right)^{\frac{1}{2}}$$

- p -norms

$$\|A\|_p = \sup_{x \neq 0} \frac{\|Ax\|_p}{\|x\|_p} = \sup_{x: \|x\|_p=1} \|Ax\|_p, \quad \|A\|_2 = \sigma_1$$

Optimal Truncation – Eckart-Young-Mirsky Theorem

Theorem (Eckart-Young-Mirsky)

For either $\|\cdot\|_2$ or $\|\cdot\|_F$,

$$\|A - A_k\| \leq \|A - B\| \text{ for all rank-}k \text{ matrices } B.$$

Moreover,

$$\|A - A_k\| = \begin{cases} \sigma_{k+1}, & \|\cdot\|_2 \text{ norm} \\ \left(\sum_{i=k+1}^r \sigma_i^2\right)^{\frac{1}{2}}, & \|\cdot\|_F \text{ norm.} \end{cases}$$

Image Compression with SVD

- Every image can be seen as a large matrix of pixel intensity values.

Image Compression with SVD

- Every image can be seen as a large matrix of pixel intensity values.
- Goal: To reduce the amount of data needed to store the image without a significant loss in visual quality.

Image Compression with SVD

- Every image can be seen as a large matrix of pixel intensity values.
- Goal: To reduce the amount of data needed to store the image without a significant loss in visual quality.
- The most important features of an image are captured by the largest singular values in the Σ matrix. By keeping only the top k singular values (and their corresponding vectors in U and V), we can create a lower-rank approximation of the original image that requires much less storage.

Image Compression with SVD

- Every image can be seen as a large matrix of pixel intensity values.
- Goal: To reduce the amount of data needed to store the image without a significant loss in visual quality.
- The most important features of an image are captured by the largest singular values in the Σ matrix. By keeping only the top k singular values (and their corresponding vectors in U and V), we can create a lower-rank approximation of the original image that requires much less storage.
- Low k : high compression and low image quality.
- High k : Low compression and high image quality.

Image Compression with SVD

- Every image can be seen as a large matrix of pixel intensity values.
- Goal: To reduce the amount of data needed to store the image without a significant loss in visual quality.
- The most important features of an image are captured by the largest singular values in the Σ matrix. By keeping only the top k singular values (and their corresponding vectors in U and V), we can create a lower-rank approximation of the original image that requires much less storage.
- Low k : high compression and low image quality.
- High k : Low compression and high image quality.
- MATLAB Demonstration (SVD and rSVD).

Matrix Completion

Matrix Completion

We often encounter datasets with missing entries. The goal of matrix completion is to intelligently fill in these missing values.

Matrix Completion

We often encounter datasets with missing entries. The goal of matrix completion is to intelligently fill in these missing values.

- **Problem:** Given a partially filled matrix, how can we accurately estimate the unknown entries?

Matrix Completion

We often encounter datasets with missing entries. The goal of matrix completion is to intelligently fill in these missing values.

- **Problem:** Given a partially filled matrix, how can we accurately estimate the unknown entries?
- **Examples:**
 - **Movie Ratings:** A user has only rated a few movies out of thousands.
 - **Image Corruption:** A photo has a scratch or a block of missing pixels.
 - **Sensor Data:** A sensor in a network fails, leaving gaps in the data.

Matrix Completion

We often encounter datasets with missing entries. The goal of matrix completion is to intelligently fill in these missing values.

- **Problem:** Given a partially filled matrix, how can we accurately estimate the unknown entries?
- **Examples:**
 - **Movie Ratings:** A user has only rated a few movies out of thousands.
 - **Image Corruption:** A photo has a scratch or a block of missing pixels.
 - **Sensor Data:** A sensor in a network fails, leaving gaps in the data.

Can we predict the values of the '?' cells?

$$\begin{bmatrix} 5 & ? & 1 & ? \\ ? & 2 & ? & 3 \\ 4 & ? & ? & 5 \\ ? & 1 & 2 & ? \end{bmatrix}$$

Matrix Completion

The core assumption is that the complete matrix has a simple, underlying structure, meaning it is low-rank.

Matrix Completion

The core assumption is that the complete matrix has a simple, underlying structure, meaning it is low-rank.

- **Why is Low Rank?**

- The columns (and rows) are not independent; they are combinations of a few "basis" vectors.
- Think of movie ratings: they are driven by a few latent factors like genre, actors, etc.

Matrix Completion

The core assumption is that the complete matrix has a simple, underlying structure, meaning it is low-rank.

- **Why is Low Rank?**

- The columns (and rows) are not independent; they are combinations of a few "basis" vectors.
- Think of movie ratings: they are driven by a few latent factors like genre, actors, etc.

- **How does SVD help?**

- SVD finds the best low-rank approximation of any matrix.
- It decomposes a matrix A into:

$$A = USV^T$$

- By keeping only the top k singular values in S and setting others to zero (S_k), we get the best rank- k approximation: $A_k = US_kV^T$.

Matrix Completion

We can't apply SVD directly to a matrix with holes. So, we:

Matrix Completion

We can't apply SVD directly to a matrix with holes. So, we:

- 1 **Initialize:** Make an initial guess for the missing entries. A simple choice is to fill them with the average of the known values, or just zero.

Matrix Completion

We can't apply SVD directly to a matrix with holes. So, we:

- 1 **Initialize:** Make an initial guess for the missing entries. A simple choice is to fill them with the average of the known values, or just zero.
- 2 **Approximate (The SVD Step):** Take this "filled" matrix and compute its SVD. Create a **low-rank approximation** by keeping only the top k singular values.

Matrix Completion

We can't apply SVD directly to a matrix with holes. So, we:

- 1 **Initialize:** Make an initial guess for the missing entries. A simple choice is to fill them with the average of the known values, or just zero.
- 2 **Approximate (The SVD Step):** Take this "filled" matrix and compute its SVD. Create a **low-rank approximation** by keeping only the top k singular values.
- 3 **Update:** Use the values from this new low-rank matrix to update your guess for the missing entries. Crucially, leave the original, known values untouched.

Matrix Completion

We can't apply SVD directly to a matrix with holes. So, we:

- 1 **Initialize:** Make an initial guess for the missing entries. A simple choice is to fill them with the average of the known values, or just zero.
- 2 **Approximate (The SVD Step):** Take this "filled" matrix and compute its SVD. Create a **low-rank approximation** by keeping only the top k singular values.
- 3 **Update:** Use the values from this new low-rank matrix to update your guess for the missing entries. Crucially, leave the original, known values untouched.
- 4 **Repeat:** Go back to Step 2 and repeat. The guesses for the missing values will converge.

Matrix Completion

Recommender Systems.

Matrix Completion

Recommender Systems.

- **Matrix:** Rows are users, columns are movies, and entries are ratings (1-5 stars).
- **Problem:** This matrix is extremely sparse—most users have rated only a tiny fraction of the available movies.
- **Goal:** Predict how a user would rate movies they haven't seen to provide personalized recommendations.

Matrix Completion

Recommender Systems.

- **Matrix:** Rows are users, columns are movies, and entries are ratings (1-5 stars).
- **Problem:** This matrix is extremely sparse—most users have rated only a tiny fraction of the available movies.
- **Goal:** Predict how a user would rate movies they haven't seen to provide personalized recommendations.

						
 a	5		1	1		2
 b		2		4		4
 c	4	5		1	1	2
 d			3	5	2	
 e	2		1		4	4

Matrix Completion

Recovering corrupted or missing parts of an image.

Matrix Completion

Recovering corrupted or missing parts of an image.

- **Matrix:** An image is just a matrix of pixel values (or three matrices for R, G, B channels).
- **Problem:** A block of pixels is missing due to a scratch, text overlay, or digital removal.
- **Goal:** Recover corrupted or missing parts of an image.

Matrix Completion

Recovering corrupted or missing parts of an image.

- **Matrix:** An image is just a matrix of pixel values (or three matrices for R, G, B channels).
- **Problem:** A block of pixels is missing due to a scratch, text overlay, or digital removal.
- **Goal:** Recover corrupted or missing parts of an image.
- MATLAB Demonstration.

Linear Regression with SVD

Linear Regression with SVD

Goal: Model a dependent variable y as a linear combination of independent variables (features) in a data matrix X .

$$y = X\beta + \epsilon$$

Linear Regression with SVD

Goal: Model a dependent variable y as a linear combination of independent variables (features) in a data matrix X .

$$y = X\beta + \epsilon$$

- y : vector of observed outcomes ($m \times 1$).
- X : design matrix of features ($m \times n$, where each row is an observation and each column is a feature).
- β : vector of unknown regression coefficients ($n \times 1$).
- ϵ : vector of random errors.

Linear Regression with SVD

Goal: Model a dependent variable y as a linear combination of independent variables (features) in a data matrix X .

$$y = X\beta + \epsilon$$

- y : vector of observed outcomes ($m \times 1$).
- X : design matrix of features ($m \times n$, where each row is an observation and each column is a feature).
- β : vector of unknown regression coefficients ($n \times 1$).
- ϵ : vector of random errors.

Objective: Find the coefficient vector β that minimizes the sum of squared errors:

$$\min_{\beta} \|y - X\beta\|_2^2$$

Linear Regression with SVD

Goal: Model a dependent variable y as a linear combination of independent variables (features) in a data matrix X .

$$y = X\beta + \epsilon$$

- y : vector of observed outcomes ($m \times 1$).
- X : design matrix of features ($m \times n$, where each row is an observation and each column is a feature).
- β : vector of unknown regression coefficients ($n \times 1$).
- ϵ : vector of random errors.

Objective: Find the coefficient vector β that minimizes the sum of squared errors:

$$\min_{\beta} \|y - X\beta\|_2^2$$

Standard Solution: $\beta = (X^T X)^{-1} X^T y$ ($X^T X$ is invertible).

Linear Regression with SVD

Challenges: Multicollinearity. This occurs when two or more feature columns in the data matrix X are highly correlated. This means they are nearly linearly dependent.

Linear Regression with SVD

Challenges: Multicollinearity. This occurs when two or more feature columns in the data matrix X are highly correlated. This means they are nearly linearly dependent.

- If columns of X are linearly dependent, the matrix $X^T X$ becomes singular (not invertible), and the normal equation solution $\beta = (X^T X)^{-1} X^T y$ cannot be computed.

Linear Regression with SVD

Challenges: Multicollinearity. This occurs when two or more feature columns in the data matrix X are highly correlated. This means they are nearly linearly dependent.

- If columns of X are linearly dependent, the matrix $X^T X$ becomes singular (not invertible), and the normal equation solution $\beta = (X^T X)^{-1} X^T y$ cannot be computed.
- If columns are nearly linearly dependent (high multicollinearity), $X^T X$ is ill-conditioned, leading to the numerical instability.

Linear Regression with SVD

SVD provides a robust and stable way to solve linear regression, even in the presence of multicollinearity.

Linear Regression with SVD

SVD provides a robust and stable way to solve linear regression, even in the presence of multicollinearity.

- The Moore-Penrose Pseudoinverse (X^+): any matrix; best-fit solution for $X\beta = y$.

Linear Regression with SVD

SVD provides a robust and stable way to solve linear regression, even in the presence of multicollinearity.

- The Moore-Penrose Pseudoinverse (X^+): any matrix; best-fit solution for $X\beta = y$.
- Computing the Pseudoinverse with SVD:

$$X = U\Sigma V^T \Rightarrow X^+ = V\Sigma^+ U^T$$

Linear Regression with SVD

SVD provides a robust and stable way to solve linear regression, even in the presence of multicollinearity.

- The Moore-Penrose Pseudoinverse (X^+): any matrix; best-fit solution for $X\beta = y$.
- Computing the Pseudoinverse with SVD:

$$X = U\Sigma V^T \Rightarrow X^+ = V\Sigma^+ U^T$$

- The regression coefficients are found by:

$$\beta = X^+ y = V\Sigma^+ U^T y$$

This solution minimizes the least-squares error $\|y - X\beta\|_2^2$.

Linear Regression with SVD

MATLAB Demonstration.

$$\begin{bmatrix} | & | & \cdots & | \\ 1 & x_1 & \cdots & x_n \\ | & | & \cdots & | \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_n \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

Summary of What You've Learned Today

- Properties of SVD, U , Σ , V .
- Computation of SVD: Conventional method and randomized SVD.
- Geometric Interpretation of SVD
- Applications
 - Matrix Approximation: Image Compression.
 - Matrix Completion: Image Inpainting.
 - Linear Regression.

Summary of What You've Learned Today

- Properties of SVD, U , Σ , V .
- Computation of SVD: Conventional method and randomized SVD.
- Geometric Interpretation of SVD
- Applications
 - Matrix Approximation: Image Compression.
 - Matrix Completion: Image Inpainting.
 - Linear Regression.

See you on Wednesday (7/2) for Lecture 3 on
Singular Value Decomposition (PCA, POD, LSA).

Principal Component Analysis (PCA)

Principal Component Analysis (PCA)

- PCA is a dimensionality reduction technique used to transform a dataset with many variables into a smaller set of new variables, called principal components (PCs), while retaining most of the original information (variance).

Principal Component Analysis (PCA)

- PCA is a dimensionality reduction technique used to transform a dataset with many variables into a smaller set of new variables, called principal components (PCs), while retaining most of the original information (variance).
- Goals of PCA:

Principal Component Analysis (PCA)

- PCA is a dimensionality reduction technique used to transform a dataset with many variables into a smaller set of new variables, called principal components (PCs), while retaining most of the original information (variance).
- Goals of PCA:
 - Reduce dimensionality for easier visualization or to combat the “curse of dimensionality”.
 - Identify key patterns and correlations in data.
 - Remove redundant features.
 - Prepare data for machine learning algorithms.

Principal Component Analysis (PCA)

- PCA is a dimensionality reduction technique used to transform a dataset with many variables into a smaller set of new variables, called principal components (PCs), while retaining most of the original information (variance).
- Goals of PCA:
 - Reduce dimensionality for easier visualization or to combat the “curse of dimensionality”.
 - Identify key patterns and correlations in data.
 - Remove redundant features.
 - Prepare data for machine learning algorithms.
- PCs are orthogonal (uncorrelated).

Principal Component Analysis (PCA)

- PCA is a dimensionality reduction technique used to transform a dataset with many variables into a smaller set of new variables, called principal components (PCs), while retaining most of the original information (variance).
- Goals of PCA:
 - Reduce dimensionality for easier visualization or to combat the “curse of dimensionality”.
 - Identify key patterns and correlations in data.
 - Remove redundant features.
 - Prepare data for machine learning algorithms.
- PCs are orthogonal (uncorrelated).
- PCs are ordered such that PC1 explains the most variance, PC2 the second most (and is orthogonal to PC1), and so on.

Principal Component Analysis (PCA)

SVD of the mean-centered data matrix $A \in \mathbb{R}^{m \times n}$ (m features, n observations).

Principal Component Analysis (PCA)

SVD of the mean-centered data matrix $A \in \mathbb{R}^{m \times n}$ (m features, n observations).

- 1 Mean-Center the Data: For each feature, subtract its mean. $A \rightarrow A_c$.
If applied to uncentered data, SVD would find directions that might be influenced by the mean itself rather than the variance structure.

Principal Component Analysis (PCA)

SVD of the mean-centered data matrix $A \in \mathbb{R}^{m \times n}$ (m features, n observations).

- 1 Mean-Center the Data: For each feature, subtract its mean. $A \rightarrow A_c$.
If applied to uncentered data, SVD would find directions that might be influenced by the mean itself rather than the variance structure.
- 2 Perform SVD on A_c : $A_c = U\Sigma V^T$.

Principal Component Analysis (PCA)

SVD of the mean-centered data matrix $A \in \mathbb{R}^{m \times n}$ (m features, n observations).

- 1 Mean-Center the Data: For each feature, subtract its mean. $A \rightarrow A_c$. If applied to uncentered data, SVD would find directions that might be influenced by the mean itself rather than the variance structure.
- 2 Perform SVD on A_c : $A_c = U\Sigma V^T$.
- 3 The columns of U are the directions of maximum variance (principal components). $PC_i = i$ -th column of U . Singular values σ_i is proportional to standard derivation along each component.

Principal Component Analysis (PCA)

SVD of the mean-centered data matrix $A \in \mathbb{R}^{m \times n}$ (m features, n observations).

- 1 Mean-Center the Data: For each feature, subtract its mean. $A \rightarrow A_c$. If applied to uncentered data, SVD would find directions that might be influenced by the mean itself rather than the variance structure.
- 2 Perform SVD on A_c : $A_c = U\Sigma V^T$.
- 3 The columns of U are the directions of maximum variance (principal components). $PC_i = i$ -th column of U . Singular values σ_i is proportional to standard derivation along each component.
- 4 Projection onto lower-dimensional subspace: $P = U^T A_c \in \mathbb{R}^{r \times n}$.

Principal Component Analysis (PCA)

MATLAB Demonstration.

Eigenfaces

Eigenfaces

Eigenfaces are the principal components of a dataset of face images.

Eigenfaces

Eigenfaces are the principal components of a dataset of face images.

- Each face image is reshaped into a column vector.

Eigenfaces

Eigenfaces are the principal components of a dataset of face images.

- Each face image is reshaped into a column vector.
- The set of all such vectors forms a data matrix $A \in \mathbb{R}^{m \times n}$, where:
 - m : number of pixels (e.g., 10000 for 100×100 images).
 - n : number of images in the dataset.

Eigenfaces

Eigenfaces are the principal components of a dataset of face images.

- Each face image is reshaped into a column vector.
- The set of all such vectors forms a data matrix $A \in \mathbb{R}^{m \times n}$, where:
 - m : number of pixels (e.g., 10000 for 100×100 images).
 - n : number of images in the dataset.
- PCA reveals the main modes of variation across faces.

Eigenfaces

Eigenfaces are the principal components of a dataset of face images.

- Each face image is reshaped into a column vector.
- The set of all such vectors forms a data matrix $A \in \mathbb{R}^{m \times n}$, where:
 - m : number of pixels (e.g., 10000 for 100×100 images).
 - n : number of images in the dataset.
- PCA reveals the main modes of variation across faces.
- The left singular vectors U from SVD are reshaped back to images.

Eigenfaces

Eigenfaces are the principal components of a dataset of face images.

- Each face image is reshaped into a column vector.
- The set of all such vectors forms a data matrix $A \in \mathbb{R}^{m \times n}$, where:
 - m : number of pixels (e.g., 10000 for 100×100 images).
 - n : number of images in the dataset.
- PCA reveals the main modes of variation across faces.
- The left singular vectors U from SVD are reshaped back to images.

$$X = U\Sigma V^T \Rightarrow \text{Eigenfaces} = U_{:,1:k}$$

Eigenfaces

Any Face as Linear Combination of Eigenfaces

Any centered face vector x can be approximated as:

$$x \approx \sum_{i=1}^k \alpha_i u_i = U_k \alpha$$

Eigenfaces

Any Face as Linear Combination of Eigenfaces

Any centered face vector x can be approximated as:

$$x \approx \sum_{i=1}^k \alpha_i u_i = U_k \alpha$$

where:

- u_i : the i -th eigenface.
- $\alpha = U_k^\top x$: projection coefficients.

Eigenfaces

Any Face as Linear Combination of Eigenfaces

Any centered face vector x can be approximated as:

$$x \approx \sum_{i=1}^k \alpha_i u_i = U_k \alpha$$

where:

- u_i : the i -th eigenface.
- $\alpha = U_k^\top x$: projection coefficients.

We can reconstruct the original face using only the top k eigenfaces:

$$x_{\text{recon}} = U_k U_k^\top x$$

Eigenfaces

Classification

Idea: Project each face image onto a low-dimensional subspace spanned by eigenfaces.

Eigenfaces

Classification

Idea: Project each face image onto a low-dimensional subspace spanned by eigenfaces.

- Each image becomes a vector of coefficients $\alpha \in \mathbb{R}^k$.
- Use classifiers (e.g., k-means, k-NN, SVM) on α instead of original image.
- Efficient and robust to noise, simple.

Proper Orthogonal Decomposition (POD)

Proper Orthogonal Decomposition

Proper Orthogonal Decomposition (POD)

Proper Orthogonal Decomposition

- Goal: To decompose a complex, high-dimensional dataset into a set of dominant, energy-ranked spatial patterns called POD modes and their corresponding temporal evolution.

Proper Orthogonal Decomposition (POD)

Proper Orthogonal Decomposition

- Goal: To decompose a complex, high-dimensional dataset into a set of dominant, energy-ranked spatial patterns called POD modes and their corresponding temporal evolution.
- POD finds the most efficient (optimal) orthogonal basis to represent the data. This means that a truncated set of POD modes captures more energy (variance) of the system than any other basis of the same size.

Proper Orthogonal Decomposition (POD)

Proper Orthogonal Decomposition

- Goal: To decompose a complex, high-dimensional dataset into a set of dominant, energy-ranked spatial patterns called POD modes and their corresponding temporal evolution.
- POD finds the most efficient (optimal) orthogonal basis to represent the data. This means that a truncated set of POD modes captures more energy (variance) of the system than any other basis of the same size.
- The decomposition takes the form:

$$X(x, t) \approx \sum_{k=1}^r a_k(t) \phi_k(x)$$

Proper Orthogonal Decomposition (POD)

Proper Orthogonal Decomposition

- Goal: To decompose a complex, high-dimensional dataset into a set of dominant, energy-ranked spatial patterns called POD modes and their corresponding temporal evolution.
- POD finds the most efficient (optimal) orthogonal basis to represent the data. This means that a truncated set of POD modes captures more energy (variance) of the system than any other basis of the same size.
- The decomposition takes the form:

$$X(x, t) \approx \sum_{k=1}^r a_k(t) \phi_k(x)$$

- $\phi_k(x)$: The k -th POD mode (spatial pattern).
- $a_k(t)$: The time coefficient for the k -th mode.

Proper Orthogonal Decomposition (POD)

The Snapshot Method: In many physical simulations (e.g., dynamics), physical quantities on a number of spatial grid points as a function of time are collected.

Proper Orthogonal Decomposition (POD)

The Snapshot Method: In many physical simulations (e.g., dynamics), physical quantities on a number of spatial grid points as a function of time are collected.

- 1 Collect Snapshots: Gather data at m different time steps. Each snapshot is a vector of all n spatial data points.

Proper Orthogonal Decomposition (POD)

The Snapshot Method: In many physical simulations (e.g., dynamics), physical quantities on a number of spatial grid points as a function of time are collected.

- 1 Collect Snapshots: Gather data at m different time steps. Each snapshot is a vector of all n spatial data points.
- 2 Form Snapshot Matrix: Arrange these snapshots as columns in a data matrix $A \in \mathbb{R}^{n \times m}$.

Proper Orthogonal Decomposition (POD)

The Snapshot Method: In many physical simulations (e.g., dynamics), physical quantities on a number of spatial grid points as a function of time are collected.

- 1 Collect Snapshots: Gather data at m different time steps. Each snapshot is a vector of all n spatial data points.
- 2 Form Snapshot Matrix: Arrange these snapshots as columns in a data matrix $A \in \mathbb{R}^{n \times m}$.
- 3 Perform SVD: $A = U\Sigma V^T$.

Proper Orthogonal Decomposition (POD)

The Snapshot Method: In many physical simulations (e.g., dynamics), physical quantities on a number of spatial grid points as a function of time are collected.

- ① **Collect Snapshots:** Gather data at m different time steps. Each snapshot is a vector of all n spatial data points.
- ② **Form Snapshot Matrix:** Arrange these snapshots as columns in a data matrix $A \in \mathbb{R}^{n \times m}$.
- ③ **Perform SVD:** $A = U \Sigma V^T$.
- ④ **Interpret the POD Components:**
 - U (Left Singular Vectors): The columns of U are the POD modes (spatial patterns). Each column is an orthonormal basis vector representing a dominant structure.
 - Σ (Singular Values): The squared singular values (σ_i^2) represent the “energy” of each mode. A rapid decay indicates the dynamics are dominated by a few modes.
 - V (Right Singular Vectors): The columns of V represent the temporal coefficients or evolution of each corresponding spatial mode.

POD on Spring-Mass Networks

MATLAB Demonstration

POD on Spring-Mass Networks

MATLAB Demonstration

- **Data Collection:**
- A 1D chain of masses connected by random springs.
- Dynamics are governed by Newton's second law:

$$M\ddot{u}(t) + Ku(t) = F(t).$$

POD on Spring-Mass Networks

MATLAB Demonstration

- **Data Collection:**
- A 1D chain of masses connected by random springs.
- Dynamics are governed by Newton's second law:

$$M\ddot{u}(t) + Ku(t) = F(t).$$

- Fix a node.
- Apply a short impulse at another node to excite the system.

POD on Spring-Mass Networks

MATLAB Demonstration

- **POD:**
- Collect displacement snapshots: $A = [u(t_1), u(t_2), \dots]$.
- Apply Singular Value Decomposition:

$$A = U\Sigma V^T.$$

POD on Spring-Mass Networks

MATLAB Demonstration

- **POD:**
- Collect displacement snapshots: $A = [u(t_1), u(t_2), \dots]$.
- Apply Singular Value Decomposition:

$$A = U\Sigma V^T.$$

- Extract:
 - Spatial modes: $\Phi = U(:, 1:r)$.
 - Time coefficients: $a(t) = \Sigma_r V_r^T$.

Latent Semantic Analysis (LSA)

What is LSA?

Latent Semantic Analysis (LSA)

What is LSA?

- LSA is a technique to analyze relationships between documents and the terms they contain.

Latent Semantic Analysis (LSA)

What is LSA?

- LSA is a technique to analyze relationships between documents and the terms they contain.
- Its goal is to uncover the **latent semantic structure** of text.

Latent Semantic Analysis (LSA)

What is LSA?

- LSA is a technique to analyze relationships between documents and the terms they contain.
- Its goal is to uncover the **latent semantic structure** of text.
- Instead of matching keywords, it understands text based on underlying **concepts** and **topics**.

Latent Semantic Analysis (LSA)

What is LSA?

- LSA is a technique to analyze relationships between documents and the terms they contain.
- Its goal is to uncover the **latent semantic structure** of text.
- Instead of matching keywords, it understands text based on underlying **concepts** and **topics**.
- This helps overcome the limitations of simple keyword matching.

Latent Semantic Analysis (LSA)

Keywords Aren't Meanings: Traditional search methods struggle with two common language issues:

Latent Semantic Analysis (LSA)

Keywords Aren't Meanings: Traditional search methods struggle with two common language issues:

1. Synonymy

Different words can have the same meaning.

- A search for “**auto**” won’t find a document that only uses the word “**car**”.

Latent Semantic Analysis (LSA)

Keywords Aren't Meanings: Traditional search methods struggle with two common language issues:

1. Synonymy

Different words can have the same meaning.

- A search for “**auto**” won’t find a document that only uses the word “**car**”.

2. Polysemy

The same word can have multiple meanings.

- A search for “**jaguar**” could return documents about the animal instead of the car brand.

Latent Semantic Analysis (LSA)

The Core Idea: The Term-Document Matrix

LSA starts by creating a matrix where:

- Each **row** corresponds to a unique **term** (word).
- Each **column** corresponds to a **document**.
- Each **cell** contains a count of how many times a term appears in a document.

Term	Doc 1	Doc 2	Doc 3	Doc 4
nasa	1	1	0	0
mars	1	0	0	0
space	1	1	0	0
data	0	0	1	1
python	0	0	1	0

Latent Semantic Analysis (LSA)

LSA uses SVD to decompose the term-document matrix (A) and then reduces its dimensions to find k concepts.

$$A \approx U_k \cdot S_k \cdot V_k^T$$

- U_k Term-Concept Matrix: Shows how each term relates to a concept.
- S_k Concept Strength Matrix: Shows the importance of each concept.
- V_k Document-Concept Matrix: Shows how each document relates to a concept.

Latent Semantic Analysis (LSA)

LSA uses SVD to decompose the term-document matrix (A) and then reduces its dimensions to find k concepts.

$$A \approx U_k \cdot S_k \cdot V_k^T$$

- U_k Term-Concept Matrix: Shows how each term relates to a concept.
- S_k Concept Strength Matrix: Shows the importance of each concept.
- V_k Document-Concept Matrix: Shows how each document relates to a concept.

This reduction forces the model to generalize, merging synonyms and clarifying ambiguous terms.

Latent Semantic Analysis (LSA)

After applying SVD, we have a low-dimensional “concept space”

- Find conceptually similar documents, even if they don't share keywords.
- Find semantically related terms.
- Compare a search query against documents in the new space.

Latent Semantic Analysis (LSA)

MATLAB Demonstration.

Latent Semantic Analysis (LSA)

MATLAB Demonstration.

Example

1. NASA sends a rover to explore the surface of Mars
2. The James Webb space telescope captures images of distant stars
3. Machine learning in Python makes data analysis simple
4. Data scientists develop new algorithms for data mining

Latent Semantic Analysis (LSA)

MATLAB Demonstration.

Example

1. NASA sends a rover to explore the surface of Mars
2. The James Webb space telescope captures images of distant stars
3. Machine learning in Python makes data analysis simple
4. Data scientists develop new algorithms for data mining

Two concepts: space exploration, data science

Summary of What You've Learned Today

- Principal Component Analysis (PCA): High Dimensional Data.
- Proper Orthogonal Decomposition (POD): Time Series.
- Latent Semantic Analysis (LSA): Natural Language Processing.

Summary of What You've Learned Today

- Principal Component Analysis (PCA): High Dimensional Data.
- Proper Orthogonal Decomposition (POD): Time Series.
- Latent Semantic Analysis (LSA): Natural Language Processing.

See you on Monday (7/7) for Lecture 4 on
Fourier Analysis.